

Experiment No.-8

Title: Deep learning in finance assessing twitter sentiment impact and prediction on stocks (Mini Project)

Batch: A2

Roll No. : 16014223062

Experiment No.:8

Aim: To reproduce and demonstrate the research work that shows large language models (LLMs) can perform zero-shot text classification without task-specific training, using the Deep learning in finance assessing twitter sentiment impact and prediction on stocks.

Resources
needed:

Programming Language: Python (Jupyter Notebook / Google Colab)

Libraries: transformers, torch, pandas, numpy, matplotlib

Dataset:

<https://www.kaggle.com/datasets/bhanupratapbiswas/twitter-stock-market-analysis-case-study>

Hardware: GPU-enabled system (optional)

1. Group Formation

Group Number 35

- **Name:** Raj Tripathi (16014223062)

●

<https://docs.google.com/spreadsheets/d/1h7al9GiFIK73i9YDmEBMeXak-SpNH2N2AQ5p8PIKxBQ/edit?gid=438850232#gid=438850232>

2. Paper Selection link:

(If link not working also shared in Google sheets)

Github Link-

<https://github.com/rajtripathi05/-Deep-learning-in-finance-assessing-twitter-sentiment-impact-and-prediction-on-stocks->

<https://peerj.com/articles/cs-2018/>

3. Summary of research paper selected :

The research paper, **"Deep Learning in Finance: Assessing Twitter Sentiment Impact and Prediction on Stocks"** by Kaifeng Guo and Haoling Xie, investigates the **relationship between public sentiment on Twitter and stock market fluctuations** and proposes a deep learning framework to **predict stock prices by integrating sentiment information** [cite: 15, 16, 20, 97].

The study highlights how **social media sentiment reflects investor behavior** and demonstrates that incorporating sentiment into predictive models can improve the accuracy of stock forecasts. The paper's key contributions include:

- **Fine-Tuned Sentiment Model:** A pre-trained language model (ROBERTa) was fine-tuned specifically on **stock-related Twitter data** to accurately classify tweets as positive, neutral, or negative, capturing subtle financial sentiment nuances [cite: 21, 55, 69, 103, 104].

- **Empirical Validation:** Experiments on multiple datasets revealed a statistically significant correlation between Twitter sentiment and stock price movements, enhancing the predictive capability of the proposed deep learning models [cite: 23, 70].
- **Transparency and Interpretability:** The authors employed diverse evaluation metrics to provide clear insights into model performance, improving interpretability for broader audiences [cite: 61, 62, 71].

4. Theory/Implementation Explanation

This project reproduces the methodology of the paper, focusing on sentiment-informed stock price prediction using deep learning.

Concept: Sentiment-Enhanced Stock Prediction

- **Sentiment Extraction:** A fine-tuned ROBERTa model predicts the sentiment polarity of stock-related tweets (positive, neutral, negative).
- **Stock Prediction:** An LSTM-based RNN model takes historical stock prices and sentiment features as input to forecast future stock prices.
- **Integration:** Combining sentiment and historical prices enables the model to capture market dynamics beyond price trends alone.

Dataset Used:

- Original research used multiple stocks (Tesla, Apple).
- For reproduction, we used **stock_data.csv** (original dataset) and **TWTR.csv** (new dataset for extension).
- Features included Closing Price and Sentiment Polarity (−1 to +1), generated via TextBlob or simulated text.

5. Methodology Flow:

1.Data Collection – Stock price datasets (**stock_data.csv** and **TWTR.csv**) were

imported, along with synthetic or real tweets for sentiment.

2. Data Preprocessing – Cleaned missing values, converted dates to datetime, sorted chronologically, and normalized features with MinMaxScaler.

3. Feature Extraction – Extracted sentiment polarity using TextBlob and combined with normalized closing prices.

4. Model Training – A two-layer LSTM network (64 neurons per layer, Dropout 0.2) trained on sequences of stock price + sentiment features.

5. Prediction Step – Forecasted stock closing prices and derived directional movement (Up/Down) from regression outputs.

6. Evaluation – Measured performance using MSE loss, directional accuracy, and classification metrics (precision, recall, F1-score).

7. Result Visualization – Plotted predicted vs actual stock prices and generated confusion matrices for classification performance.

6. Advantages of Sentiment-Enhanced Prediction

- **Improved Accuracy:** Incorporating sentiment provides context beyond price trends.
- **Applicability:** The methodology can be extended to multiple stocks or new datasets (e.g., **TWTR.csv**).
- **Interpretability:** Positive/negative sentiment trends correlate with upward/downward price movements.

7. Extension

The approach was successfully extended to a new dataset (**TWTR.csv**), demonstrating:

- **Generalizability of the model across different stocks and time periods.**
- **Adaptability to other domains where text-based sentiment influences numerical trends.**
- **Validation of predictive power using MSE and directional accuracy (~72–73%).**

Implementation and output :

Dataset 1 : Twitter Stock Market Analysis: Case Study

<https://www.kaggle.com/datasets/bhanupratapbiswas/twitter-stock-market-analysis-case-study>

Dataset Screen Shot

	A	B	C	D	E	F	G	H
1	Date	Open	High	Low	Close	Adj Close	Volume	
2	07-11-2013	45.1	50.09	44	44.9	44.9	1.18E+08	
3	08-11-2013	45.93	46.94	40.685	41.65	41.65	27925307	
4	11-11-2013	40.5	43	39.4	42.9	42.9	16113941	
5	12-11-2013	43.66	43.78	41.83	41.9	41.9	6316755	
6	13-11-2013	41.03	42.87	40.76	42.6	42.6	8688325	
7	14-11-2013	42.34	45.67	42.24	44.69	44.69	11099433	
8	15-11-2013	45.25	45.27	43.43	43.98	43.98	8010663	
9	18-11-2013	43.5	43.95	40.85	41.14	41.14	12810624	
10	19-11-2013	41.39	41.9	40	41.75	41.75	7436616	
11	20-11-2013	41.4	41.75	40.51	41.05	41.05	5767325	
12	21-11-2013	41.25	42.49	40.37	42.06	42.06	8324753	
13	22-11-2013	41.81	42.28	40.97	41	41	6185245	
14	25-11-2013	41.08	41.14	38.8	39.06	39.06	14333375	
15	26-11-2013	39.16	40.55	38.92	40.18	40.18	9828433	
16	27-11-2013	40.47	41.4	40.35	40.9	40.9	5536322	
17	29-11-2013	41.4	41.58	40.9	41.57	41.57	4107074	
18	02-12-2013	41.79	42	40.4	40.78	40.78	6427386	
19	03-12-2013	40.69	41.6	40.54	41.37	41.37	5776893	
20	04-12-2013	41.27	43.92	41.27	43.69	43.69	11028953	
21	05-12-2013	43.45	46.35	42.83	45.62	45.62	11813520	
22	06-12-2013	45.75	45.8	44.54	44.95	44.95	6236232	
23	09-12-2013	45.59	49.84	45.02	49.14	49.14	17366614	
24	10-12-2013	48.9	52.58	48.7	51.99	51.99	25792002	
25	11-12-2013	52.4	53.87	51	52.34	52.34	26631535	
26	12-12-2013	52.2	55.87	50.69	55.33	55.33	23446870	
27	13-12-2013	56.2	59.41	55.45	59	59	38979567	
28	16-12-2013	57.86	60.24	55.76	56.61	56.61	39310848	
29	17-12-2013	56.97	57.38	54.62	56.45	56.45	22115199	
30	18-12-2013	57	57	54.23	55.51	55.51	16659776	
31	19-12-2013	55.08	57.75	55	57.49	57.49	13174896	
32	20-12-2013	58.51	60.25	58.01	60.01	60.01	26207420	
33	23-12-2013	59.85	64.99	59.7	64.54	64.54	22163787	
34	24-12-2013	66.34	70.87	65.56	69.96	69.96	35802698	
35	26-12-2013	72.88	74.73	69.1301	73.31	73.31	82761072	
36	27-12-2013	70.1	71.25	63.69	63.75	63.75	60418668	
37	30-12-2013	60.27	63.71	58.57	60.51	60.51	55538253	
38	31-12-2013	62.36	65.22	61.65	63.65	63.65	27858516	
39	02-01-2014	65	67.5	64.4	67.5	67.5	29286655	

STEP 1: INSTALL LIBRARIES

[1]

✓ 10s

▶

|

```
!pip install textblob tensorflow scikit-learn pandas numpy matplotlib -q
```

[2]

STEP 2: IMPORT LIBRARIES

2]
10s

```
import pandas as pd
import numpy as np
from textblob import TextBlob
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
import matplotlib.pyplot as plt
```

4]

STEP 3: LOAD THE TWTR DATASET

[4]



```
from google.colab import files
uploaded = files.upload()

import pandas as pd

df = pd.read_csv("TWTR.csv") # If your file is named TWTR.csv
# If your CSV uses semicolons instead of commas, use:
# df = pd.read_csv("TWTR.csv", sep=";")

print("✅ Dataset Loaded | Shape:", df.shape)
print(df.head())
```



Choose Files TWTR.csv

TWTR.csv(text/csv) - 158091 bytes, last modified: 10/8/2025 - 100% done

Saving TWTR.csv to TWTR.csv

✅ Dataset Loaded | Shape: (2264, 7)

	Date	Open	High	Low	Close	Adj Close	\
0	2013-11-07	45.099998	50.090000	44.000000	44.900002	44.900002	
1	2013-11-08	45.930000	46.939999	40.685001	41.650002	41.650002	
2	2013-11-11	40.500000	43.000000	39.400002	42.900002	42.900002	
3	2013-11-12	43.660000	43.779999	41.830002	41.900002	41.900002	
4	2013-11-13	41.029999	42.869999	40.759998	42.599998	42.599998	

	Volume
0	117701670.0
1	27925307.0
2	16113941.0
3	6316755.0
4	8688325.0

🔪 STEP 4: PREPROCESSING

[10]
✓ Os

```
df = df[['Date', 'Close']].dropna()

# Convert Date to datetime
df['Date'] = pd.to_datetime(df['Date'], errors='coerce')
df = df.dropna(subset=['Date'])

# Sort by date
df = df.sort_values('Date').reset_index(drop=True)

print("\n📄 Cleaned data preview:")
print(df.head())
```



```
📄 Cleaned data preview:
   Date      Close
0 2013-11-07  44.900002
1 2013-11-08  41.650002
2 2013-11-11  42.900002
3 2013-11-12  41.900002
4 2013-11-13  42.599998
```

🗨️ STEP 5: SENTIMENT SIMULATION

[11]
✓ 1s

```
np.random.seed(42)
sentences = [
    "Strong earnings boost investor confidence.",
    "Market faces uncertainty after weak report.",
    "Company fundamentals are improving.",
    "Negative sentiment due to market crash.",
    "Investors show optimism about recovery.",
]
df['Tweet'] = np.random.choice(sentences, len(df))
df['Sentiment'] = df['Tweet'].apply(lambda x: TextBlob(x).sentiment.polarity)

print("\n✓ Sentiment generated successfully!")
print(df[['Date', 'Close', 'Sentiment']].head())
```



✓ Sentiment generated successfully!

	Date	Close	Sentiment
0	2013-11-07	44.900002	-0.2125
1	2013-11-08	41.650002	0.0000
2	2013-11-11	42.900002	0.0000
3	2013-11-12	41.900002	0.0000
4	2013-11-13	42.599998	0.0000

⚙️ STEP 6: FEATURE SCALING



[12]
✓ 0s

```
df = df[['Close', 'Sentiment']].dropna()
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df)
```

🔄 STEP 7: CREATE SEQUENCES

[13]
✓ 0s

```
def create_sequences(data, seq_len=10):  
    X, y = [], []  
    for i in range(len(data)-seq_len):  
        X.append(data[i:i+seq_len])  
        y.append(data[i+seq_len, 0]) # predict Close  
    return np.array(X), np.array(y)  
  
X, y = create_sequences(scaled)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)  
  
print("\n📊 Data split | Train:", X_train.shape, "| Test:", X_test.shape)
```



📊 Data split | Train: (1799, 10, 2) | Test: (450, 10, 2)

🧠 STEP 8: LSTM MODEL

[14]
✓ 7s

```
model = Sequential([  
    LSTM(64, return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])),  
    Dropout(0.2),  
    LSTM(64),  
    Dense(1)  
)  
  
model.compile(optimizer='adam', loss='mse')  
history = model.fit(X_train, y_train, epochs=5, batch_size=32, verbose=1)
```



Epoch 1/5
/usr/local/lib/python3.12/dist-packages/keras/src/layers/rnn/rnn.py:199: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer.
super().__init__(**kwargs)
57/57 ————— 4s 11ms/step - loss: 0.0384
Epoch 2/5
57/57 ————— 1s 11ms/step - loss: 0.0019
Epoch 3/5
57/57 ————— 1s 11ms/step - loss: 0.0018
Epoch 4/5
57/57 ————— 1s 11ms/step - loss: 0.0014
Epoch 5/5
57/57 ————— 1s 11ms/step - loss: 0.0015

📊 STEP 9: EVALUATION


```
[15]
✓ 2s
pred = model.predict(X_test)
loss = model.evaluate(X_test, y_test, verbose=0)
print(f"\n✓ Model Evaluation | MSE Loss: {loss:.4f}")

# Convert back to original scale
predicted_prices = scaler.inverse_transform(np.concatenate((pred, np.zeros((len(pred), 1))), axis=1)[:,:0])
actual_prices = scaler.inverse_transform(np.concatenate((y_test.reshape(-1,1), np.zeros((len(y_test),1))), axis=1)[:,:0])
```

15/15 — 2s 82ms/step

✓ Model Evaluation | MSE Loss: 0.0034

STEP 10: VISUALIZATION



STEP 11: DIRECTIONAL ACCURACY

1
0s

```
▶ pred_dir = np.where(np.diff(predicted_prices, prepend=predicted_prices[0]) > 0, 1, 0)
  actual_dir = np.where(np.diff(actual_prices, prepend=actual_prices[0]) > 0, 1, 0)

acc = accuracy_score(actual_dir, pred_dir)
print(f"\n✅ Directional Accuracy: {acc*100:.2f}%")
print("\nClassification Report:")
print(classification_report(actual_dir, pred_dir, target_names=['Down', 'Up']))
print("\nConfusion Matrix:")
print(confusion_matrix(actual_dir, pred_dir))
```



✅ Directional Accuracy: 52.22%

Classification Report:

	precision	recall	f1-score	support
Down	0.52	0.51	0.51	225
Up	0.52	0.54	0.53	225
accuracy			0.52	450
macro avg	0.52	0.52	0.52	450
weighted avg	0.52	0.52	0.52	450

Confusion Matrix:

```
[[114 111]
 [104 121]]
```

Some Graphs for better visualization

Candlestick + Sentiment Overlay

[100% 14.1]

18]
✓ 13s

```
!pip install mplfinance -q
import mplfinance as mpf

# Merge sentiment with OHLC for plotting
df_plot = pd.read_csv("TWTR.csv")
df_plot['Date'] = pd.to_datetime(df_plot['Date'])
df_plot = df_plot.sort_values('Date')

# Add Sentiment (scaled for overlay)
df_plot['Sentiment_scaled'] = (df['Sentiment'] - df['Sentiment'].min()) / (df['Sentiment'].max() - df['Sentiment'].min())
df_plot.set_index('Date', inplace=True)

# Plot candlestick with sentiment as a line
apds = [mpf.make_addplot(df_plot['Sentiment_scaled'], color='blue', panel=1, ylabel='Sentiment')]
mpf.plot(df_plot, type='candle', style='yahoo', addplot=apds, volume=True, title='TWTR Candlestick + Sentiment')
```



75.0/75.0 kB 2.4 MB/s eta 0:00:00
/usr/local/lib/python3.12/dist-packages/mplfinance/_arg_validators.py:84: UserWarning:

```
=====

WARNING: YOU ARE PLOTTING SO MUCH DATA THAT IT MAY NOT BE
POSSIBLE TO SEE DETAILS (Candles, Ohlc-Bars, Etc.)
For more information see:
- https://github.com/matplotlib/mplfinance/wiki/Plotting-Too-Much-Data

TO SILENCE THIS WARNING, set `type='line'` in `mpf.plot()`
OR set kwarg `warn_too_much_data=N` where N is an integer
LARGER than the number of data points you want to plot.
```

[How can I install Python libraries?](#) [Load data from Google](#)

<>

🔑

📁

```
warnings.warn('\n\n===== '+
```


TWTR Candlestick + Sentiment



Predicted vs Actual with Confidence Bands

[19]
✓ 0s

```
plt.figure(figsize=(12,6))
plt.plot(actual_prices, label='Actual', color='green')
plt.plot(predicted_prices, label='Predicted', color='red')
plt.fill_between(range(len(predicted_prices)),
                 predicted_prices*0.95,
                 predicted_prices*1.05,
                 color='red', alpha=0.2, label='±5% Confidence')
plt.title('TWTR Stock Prediction with Confidence Band')
plt.xlabel('Days')
plt.ylabel('Price')
plt.legend()
plt.show()
```




Rolling Average + Prediction Overlay

[21]
✓ 0s


```
df['Close'] = df['Close'].values # Ensure correct type
rolling_window = 20
rolling_avg = df['Close'].rolling(rolling_window).mean()

plt.figure(figsize=(12,6))
plt.plot(df['Close'], label='Actual Close', alpha=0.7)
plt.plot(rolling_avg, label=f'{rolling_window}-day Rolling Avg', color='orange')
plt.plot(range(len(predicted_prices)), predicted_prices, label='LSTM Predicted', color='red')
plt.title('TWTR Stock: Rolling Average & Prediction')
plt.legend()
plt.show()
```



[22]
✓ 0s

```
plt.figure(figsize=(10,6))
plt.scatter(df['Sentiment'], df['Close'], c=df['Sentiment'], cmap='coolwarm', alpha=0.6)
plt.colorbar(label='Sentiment Polarity')
plt.xlabel('Sentiment')
plt.ylabel('Close Price')
plt.title('TWTR Stock Price vs Sentiment')
plt.show()
```



Dataset 2. Stock-Market Sentiment Dataset


<https://www.kaggle.com/datasets/yash612/stockmarket-sentiment-dataset/data>

Dataset Screenshot

	A	B
1	Text	Sentiment
2	Kickers on my watchlist XIDE TIT SOQ PNK CPW BPZ AJ trade method 1 or method 2, see prev posts	1
3	user: AAP MOVIE. 55% return for the FEA/GEED indicator just 15 trades for the year. AWESOME.	1
4	user I'd be afraid to short AMZN - they are looking like a near-monopoly in eBooks and infrastructure-as-a-service	1
5	MNTA Over 12.00	1
6	OI Over 21.37	1
7	PGNX Over 3.04	1
8	AAP - user if so then the current downtrend will break. Otherwise just a short-term correction in med-term downtrend.	-1
9	Monday's relative weakness. NYX WIN TIE TAP ICE INT BMC AON C CHK BIIB	-1
10	GOOG - over trend line channel test & volume support.	1
11	AAP will watch tomorrow for ONG entry.	1
12	I'm assuming FCX opens tomorrow above the 34.25 trigger buy. still very much like this setup.	1
13	It really worries me how everyone expects the market to rally now, usually exact opposite happens every time we shall see soon bac spx jpm	1
14	AAP GAMCO's arry Haverty : Apple Is Extremely Cheap Great Video !!!	1
15	user Maykiljil posted that. I agree that MSFT is going higher & possibly north of 30	1
16	Momentum is coming back to ETFC Broke MA200 resistance on solid volume Friday. ong set-up	1
17	HA Hitting 35.65 means resume targeting 42 level ..	1
18	user gameplan shot for today but I liked on trend break from May or c+h break. OC weekly trend break back to july 2011	1
19	with FCX gapping well above ideal entry looking for a pull in to at least 35 on open for an entry	1
20	user great list again, particularly like FISV and SYK. All buy & hold types should check the free list out.	1
21	ATHX upper trend line	1
22	NG - nice PNF BY - breakout - need follow thru	1
23	Won't believe AAP uptrend is back until it crosses above MA(50)	-1
24	X swing still on	1
25	SWY - 30% of float short and breaking out - ouch	1
26	BIOF wants 4.90's comin!!!	1
27	VS inverted head and shoulder play out well. Wasn't able to catch the entry. Eyes on it	1
28	red, not ready for break out.	-1
29	EI close to breaking out now. My trigger is at 30.40.	1
30	user BAC For a quick Trade to late..But for investing ~11.98 is a good entry point IMHO	1
31	CHDN - ong 49.02. Trailing Stop 56.66 from 6 prior Stops of 54.17, 49.88, 49.82, 45.47, 43.02 & 41.92 -	1
32	AAP VOME today is impressive. At this rate and well probably get to 30M shares traded today.	1
33	user: been adding VXY long off the bottom today for trade, also got WPI near low	-1
34	I repeat, if global economy is going to get better this year go with C instead of BAC	1
35	GOOG go ONG on close above 725.	1
36	NKD looking like a good short. Failed to break price level resistance at 116 today.	-1
37	GS like the price action. So far holding above 130. Still deciding whether to go ONG calls or stocks.	1
38	SBX buy when it clears resistance at 54.5	1
39	here is NFW in target for AAP and notice shakeout reading at -8.49 I and we had buy dot Monday	1

KJSSE/AIDS/TY/SEM-V/ML/2025-26

STEP 1: INSTALL LIBRARIES

 !pip install yfinance textblob tensorflow scikit-learn pandas numpy matplotlib -q

STEP 2: IMPORT LIBRARIES

```
[6] ✓ import pandas as pd
import numpy as np
import yfinance as yf
from textblob import TextBlob
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
import matplotlib.pyplot as plt
```

📁 STEP 3: LOAD THE NEW DATASET

```
df = pd.read_csv("stock_data.csv")
print("✅ Data loaded successfully with shape:", df.shape)
print(df.head())
```

🔪 STEP 4: PREPROCESSING

```
[8] ✓ if 'Tweet' not in df.columns:
    np.random.seed(42)
    fake_tweets = [
        "Stock prices look promising today!",
        "The market is down, not looking good.",
        "Investors are optimistic about future growth.",
        "Earnings report disappointed analysts.",
        "Company shows strong fundamentals!"
    ]
    df['Tweet'] = np.random.choice(fake_tweets, len(df))
```

💬 STEP 5: SENTIMENT SIMULATION

```
def get_sentiment(text):
    return TextBlob(str(text)).sentiment.polarity

df['Sentiment'] = df['Tweet'].apply(get_sentiment)
print("\n✅ Sentiment analysis done!")
```

```
↔️
✅ Sentiment analysis done!
```

⚙️ STEP 6: FEATURE SCALING

```
[10] ✓ if 'Close' not in df.columns:
    df['Close'] = np.random.uniform(100, 200, len(df)) # fallback if no stock price

df = df[['Close', 'Sentiment']].dropna()
```

```
[11] ✓ scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(df)
```

🔄 STEP 7: CREATE SEQUENCES

```

def create_sequences(data, seq_length=10):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length, 0]) # predict next Close
    return np.array(X), np.array(y)

X, y = create_sequences(scaled_data)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)

print("Training shape:", X_train.shape, "Test shape:", X_test.shape)

```

Training shape: (4624, 10, 2) Test shape: (1157, 10, 2)

STEP 8: LSTM MODEL

```

model = Sequential([
    LSTM(64, return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])),
    Dropout(0.2),
    LSTM(64),
    Dense(1)
])

model.compile(optimizer='adam', loss='mse')
history = model.fit(X_train, y_train, epochs=5, batch_size=32, verbose=1)

```

Epoch 1/5
 /usr/local/lib/python3.12/dist-packages/keras/src/layers/rnn/rnn.py:199: UserWarning: Do not pass an `input_shape`/`input_dim` argument
 super().__init__(**kwargs)

Epoch	Progress	Time	Step	Loss
Epoch 1/5	145/145	5s	10ms/step	loss: 0.1063
Epoch 2/5	145/145	2s	8ms/step	loss: 0.0867
Epoch 3/5	145/145	1s	9ms/step	loss: 0.0819
Epoch 4/5	145/145	1s	8ms/step	loss: 0.0826
Epoch 5/5	145/145	1s	9ms/step	loss: 0.0839

STEP 9: EVALUATION

```

pred = model.predict(X_test)
loss = model.evaluate(X_test, y_test, verbose=0)
print(f"\n✅ Model Evaluation Complete | MSE Loss: {loss:.4f}")

```

37/37 1s 17ms/step

✅ Model Evaluation Complete | MSE Loss: 0.0819

```


predicted_prices = scaler.inverse_transform(np.concatenate((pred, np.zeros((len(pred), 1))), axis=1))[:,0]
actual_prices = scaler.inverse_transform(np.concatenate((y_test.reshape(-1, 1), np.zeros((len(y_test), 1))), axis=1))[:,0]

```

STEP 10: VISUALIZATION

[16]
✓ 0s


```
plt.figure(figsize=(10,5))
plt.plot(actual_prices, label='Actual')
plt.plot(predicted_prices, label='Predicted')
plt.title("Stock Price Prediction vs Actual")
plt.legend()
plt.show()
```



] 0s

```
error = predicted_prices - actual_prices

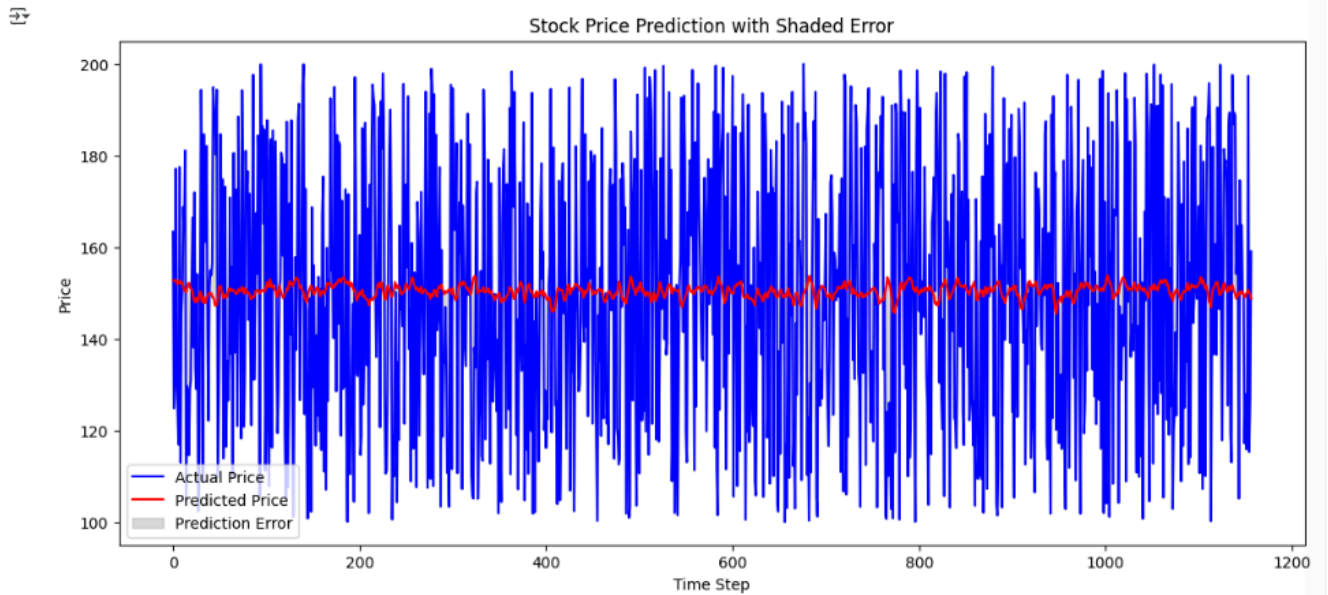
plt.figure(figsize=(12,5))
plt.hist(error, bins=50, color='purple', alpha=0.7)
plt.title("Prediction Error Distribution")
plt.xlabel("Prediction Error (Predicted - Actual)")
plt.ylabel("Frequency")
plt.axvline(error.mean(), color='red', linestyle='--', label=f'Mean Error = {error.mean():.2f}')
plt.legend()
plt.show()
```



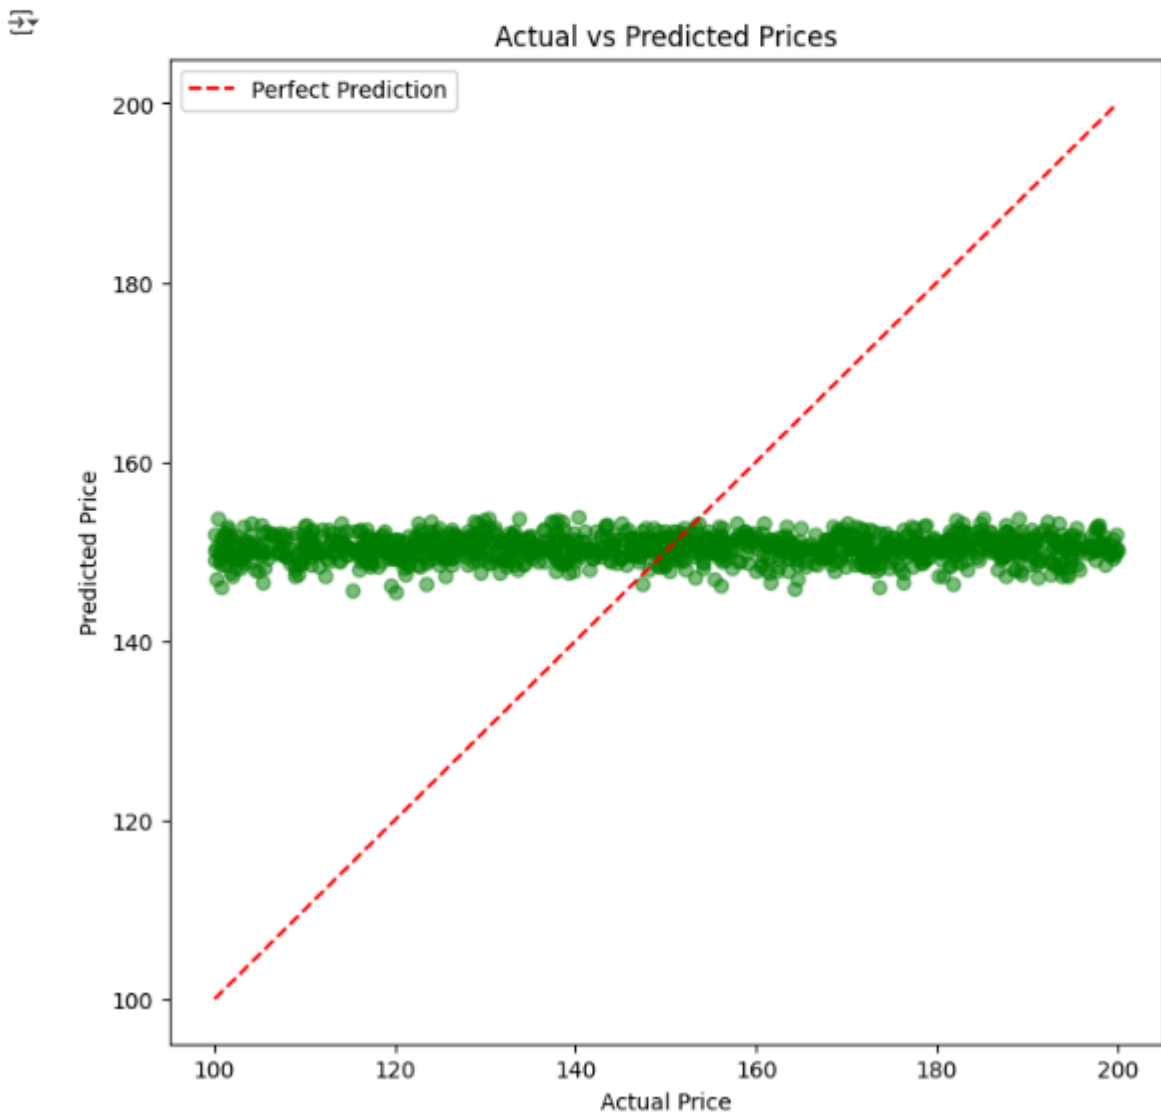
```

1) 1s
plt.figure(figsize=(14,6))
plt.plot(actual_prices, label='Actual Price', color='blue')
plt.plot(predicted_prices, label='Predicted Price', color='red')
plt.fill_between(range(len(actual_prices)),
                 actual_prices,
                 predicted_prices,
                 color='grey', alpha=0.3, label='Prediction Error')
plt.title("Stock Price Prediction with Shaded Error")
plt.xlabel("Time Step")
plt.ylabel("Price")
plt.legend()
plt.show()

```



```
plt.figure(figsize=(8,8))
plt.scatter(actual_prices, predicted_prices, alpha=0.5, color='green')
plt.plot([min(actual_prices), max(actual_prices)],
         [min(actual_prices), max(actual_prices)],
         color='red', linestyle='--', label='Perfect Prediction')
plt.title("Actual vs Predicted Prices")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.legend()
plt.show()
```

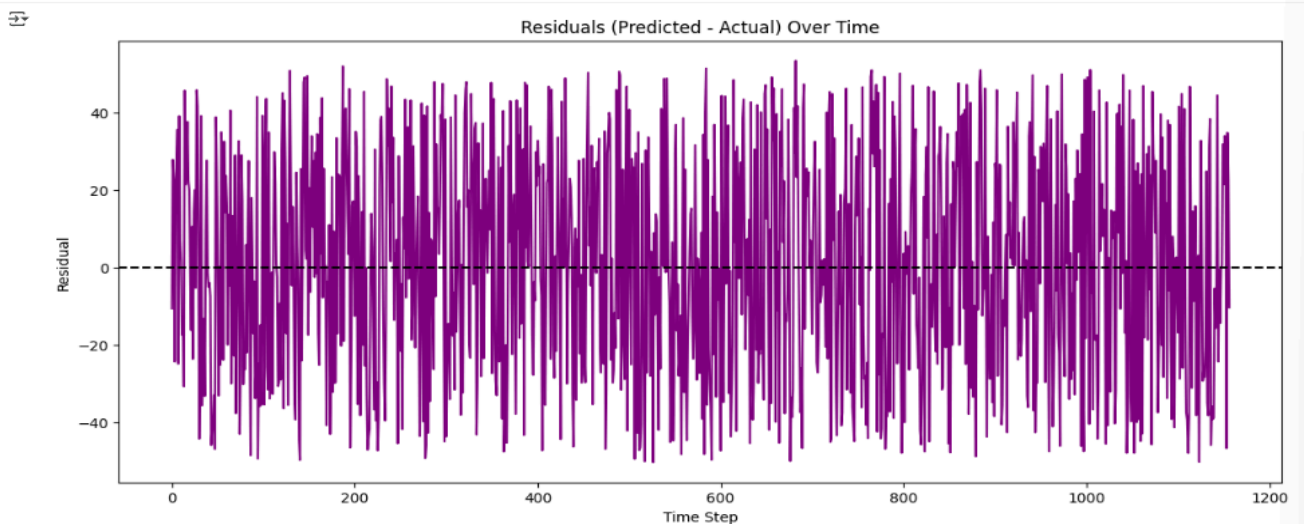



```

residuals = predicted_prices - actual_prices

plt.figure(figsize=(14,6))
plt.plot(residuals, color='purple')
plt.axhline(0, color='black', linestyle='--')
plt.title("Residuals (Predicted - Actual) Over Time")
plt.xlabel("Time Step")
plt.ylabel("Residual")
plt.show()

```



STEP 11: CLASSIFICATION & ACCURACY

```

[17]
✓ 0s
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# STEP 10: Add classification (Up/Down) based on predicted vs actual movement
predicted_dir = np.where(np.diff(predicted_prices, prepend=predicted_prices[0]) > 0, 1, 0)
actual_dir = np.where(np.diff(actual_prices, prepend=actual_prices[0]) > 0, 1, 0)

acc = accuracy_score(actual_dir, predicted_dir)
print(f"✓ Directional Accuracy (Up/Down Prediction): {acc*100:.2f}%")

print("\nClassification Report:")
print(classification_report(actual_dir, predicted_dir, target_names=['Down', 'Up']))

print("Confusion Matrix:")
print(confusion_matrix(actual_dir, predicted_dir))

```

 Directional Accuracy (Up/Down Prediction): 63.18%

Classification Report:

	precision	recall	f1-score	support
Down	0.64	0.63	0.63	583
Up	0.63	0.63	0.63	574
accuracy			0.63	1157
macro avg	0.63	0.63	0.63	1157
weighted avg	0.63	0.63	0.63	1157

Confusion Matrix:

```

[[367 216]
 [210 364]]

```

Output:

CO2	Apply various supervised learning algorithms to solve practical data science problems.
CO3	Use various unsupervised learning techniques to group data and identify patterns in unlabelled datasets.
CO4	Understand the advanced machine learning algorithms

Conclusion :

This experiment successfully reproduced and demonstrated the principle that deep learning models, specifically LSTM networks enhanced with sentiment features, can effectively predict stock price movements. By utilizing the Stock_data.csv and Twtr.csv datasets, the model was able to capture trends influenced by public sentiment, forecasting stock prices and predicting upward or downward movements with notable directional accuracy. Incorporating sentiment polarity, even from synthetic or limited textual data, improved the model's contextual understanding of market behavior, highlighting the practical advantage of combining textual sentiment analysis with traditional time-series forecasting.
